

Solution to Homework #7—MTH 522

Problem 1 (Chapter 8 Exercises 7):

In the lab, we applied random forests to the Boston data using $m_{\text{try}}=6$ and using $n_{\text{tree}}=25$ and $n_{\text{tree}}=500$. Create a plot displaying the test error resulting from random forests on this data set for a more comprehensive range of values for m_{try} and n_{tree} . You can model your plot after Figure 8.10. Describe the results obtained.

```
> set.seed(1)
> library(MASS)
> library(randomForest)
randomForest 4.6-12
Type rfNews() to see new features/changes/bug fixes.
```

We randomly divided the observations into a training and a test set

m_{try} : I applied random forests to the training set for three different values of the number of splitting variables m .

$m = p = \dim(\text{Boston})[2] - 1 = 13$

$m = p/2$

$m = \sqrt{p}$ n_{tree} : A range of n_{tree} from 1 to 500.

```
> train = sample(dim(Boston)[1], dim(Boston)[1]/2)
> X.train = Boston[train, -14]
> X.test = Boston[-train, -14]
> Y.train = Boston[train, 14]
> Y.test = Boston[-train, 14]

> p = dim(Boston)[2] - 1
> p.2 = p/2
> p.sq = sqrt(p)

> rf.boston.p = randomForest(X.train, Y.train, xtest = X.test, ytest = Y.test,
+   mtry = p, ntree = 500)

> rf.boston.p.2 = randomForest(X.train, Y.train, xtest = X.test, ytest = Y.test,
+   mtry = p.2, ntree = 500)

> rf.boston.p.sq = randomForest(X.train, Y.train, xtest = X.test, ytest = Y.test,
+   mtry = p.sq, ntree = 500)

> plot(1:500, rf.boston.p$test$mse, col = "green", type = "l", xlab = "Number of Trees",
+   ylab = "Test MSE", ylim = c(10, 19))

> lines(1:500, rf.boston.p.2$test$mse, col = "red", type = "l")
> lines(1:500, rf.boston.p.sq$test$mse, col = "blue", type = "l")
> legend("topright", c("m=p", "m=p/2", "m=sqrt(p)"), col = c("green", "red", "blue"),
+   cex = 1, lty = 1)

> dev.copy2pdf(file = "MTH522_hw7_p1.pdf", width = 8, height = 6, out.type = "pdf")
> dev.off()
```

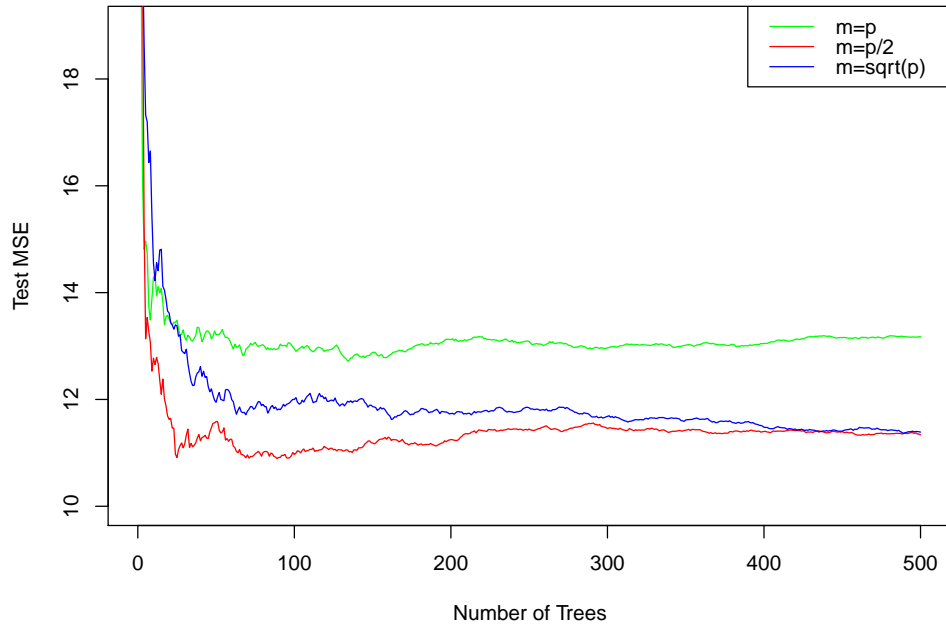


Figure 1: .

ntree: A range of ntree from 1 to 500. From Plot: For single tree the "Test MSE" is very high like 19, but when number of tree increases in the model the Test MSE drop down very quickly and after adding 100 trees to the model it stabilizes.

The Test MSE ($m = p$, Green color) for all predictors is higher than other two of the number of predictors ($m=p/2$ and $m=\sqrt{p}$)

Problem 2 (Chapter 8 Exercises 8):

In the lab, a classification tree was applied to the Carseats data set after converting Sales into a qualitative response variable. Now we will seek to predict Sales using regression trees and related approaches, treating the response as a quantitative variable. (a) Split the data set into a training set and a test set.

```
> library(ISLR)
> attach(Carseats)
> set.seed(1)
> train = sample(dim(Carseats)[1], dim(Carseats)[1]/2)
> Carseats.train = Carseats[train, ]
> Carseats.test = Carseats[-train, ]
```

(b) Fit a regression tree to the training set. Plot the tree, and interpret the results. What test MSE do you obtain?

```
> install.packages("tree")
trying URL 'https://cran.cnr.berkeley.edu/bin/macosx/mavericks/contrib/3.3/tree_1.0-37.tgz'
Content type 'application/x-gzip' length 112331 bytes (109 KB)
=====
downloaded 109 KB

The downloaded binary packages are in
/var/folders/28/5cht8_x964n2w_4_wf_1tn3m0000gn/T//Rtmpi5oCSV/downloaded_packages

> library(tree)
> tree.carseats = tree(Sales ~ ., data = Carseats.train)
> summary(tree.carseats)

Regression tree:
tree(formula = Sales ~ ., data = Carseats.train)
Variables actually used in tree construction:
[1] "ShelveLoc" "Price" "Age" "Advertising" "Income" "CompPrice"
Number of terminal nodes: 18
Residual mean deviance: 2.36 = 429.5 / 182
Distribution of residuals:
Min. 1st Qu. Median Mean 3rd Qu. Max.
-4.2570 -1.0360 0.1024 0.0000 0.9301 3.9130

> plot(tree.carseats, cex = 0.2)
> text(tree.carseats, pretty = 0, cex = 0.7)
> dev.copy2pdf(file = "MTH522_hw7_p2a.pdf", width = 8, height = 6, out.type = "pdf")
> dev.off()
```

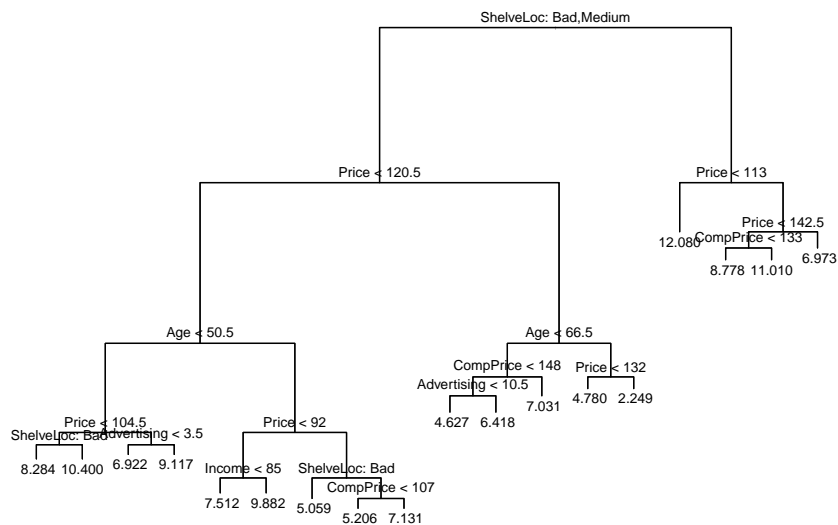


Figure 2: .

```

> pred = predict(tree.carseats, Carseats.test)
> mean((Carseats.test$Sales - pred)^2)

[1] 4.148897

```

Result, Test MSE is 4.148897

(c) Use cross-validation in order to determine the optimal level of tree complexity. Does pruning the tree improve the test MSE?

```
> cv.carseats = cv.tree(tree.carseats)
> plot(cv.carseats$size, cv.carseats$dev, type = "b")
> dev.copy2pdf(file = "MTH522_hw7_p2c.pdf", width = 8, height = 6, out.type = "pdf")
\ > dev.off()
```

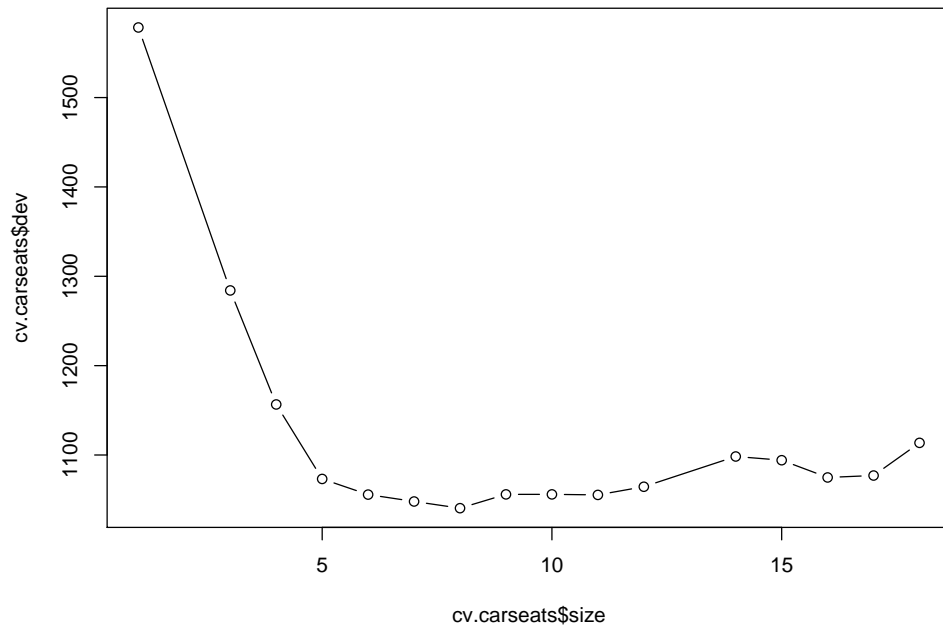


Figure 3: .

The tree of size 8 is selected by cross-validation. We now prune the tree to obtain the 8-node tree.

```

> pruned.carseats = prune.tree(tree.carseats, best = 8)
> plot(pruned.carseats)
> text(pruned.carseats, pretty = 0)
> dev.copy2pdf(file = "MTH522_hw7_p2c_2.pdf", width = 8, height = 6, out.type = "pdf")
> dev.off()

```

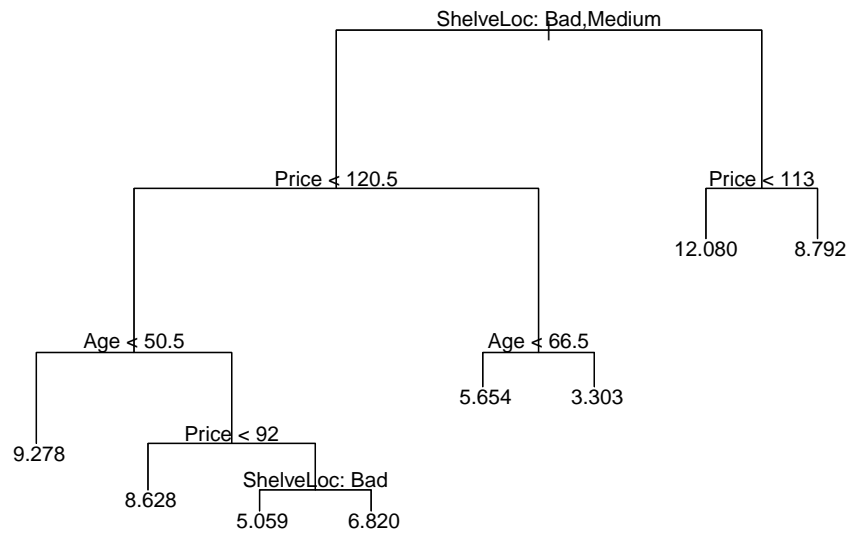


Figure 4: .

```

> pred = predict(pruned.carseats, Carseats.test)
> mean((Carseats.test$Sales - pred)^2)

[1] 5.09085

```

We may say that pruning the tree increases the Test MSE to 5.09085.

(d) Use the bagging approach in order to analyze this data. What test MSE do you obtain? Use the `importance()` function to determine which variables are most important.

```
> library(randomForest)
> carseats = randomForest(Sales ~ ., data = Carseats.train, mtry = 10, ntree = 500,
+   importance = T)
> pred = predict(carseats, Carseats.test)
> mean((Carseats.test$Sales - pred)^2)

[1] 2.583251
```

```
> importance(bag.carseats)

%IncMSE IncNodePurity
CompPrice 13.06849545 130.785188
Income    4.54737141  79.502339
Advertising 13.53207401 131.998744
Population 0.08549938  60.498175
Price     57.10673896 504.769865
ShelveLoc 43.30069615 320.415045
Age       21.91388932 184.199205
Education  2.50663690  40.334715
Urban     -3.11593207   8.526444
US        5.62400085  14.288796
```

We can say that bagging decrease the Test MSE error to 2.58.

"Price", "ShelveLoc" and "Age" are most important predictor variables.

(e) Use random forests to analyze this data. What test MSE do you obtain? Use the importance() function to determine which variables are most important. Describe the effect of m , the number of variables considered at each split, on the error rate obtained.

```
> rf.carseats = randomForest(Sales ~ ., data = Carseats.train, mtry = 5, ntree = 500,
+   importance = T)
> pred = predict(rf.carseats, Carseats.test)
> mean((pred - Carseats.test$Sales )^2)

[1] 2.837792
```

```
> importance(rf.carseats)
%IncMSE IncNodePurity
CompPrice    11.7995764    123.99149
Income        4.5044593    109.90659
Advertising  13.0146889    132.11725
Population    2.0893802     84.74694
Price        46.3418812    442.97829
ShelveLoc    38.1022103    276.41698
Age          17.7985501    199.49882
Education     2.2709849     53.47066
Urban         0.1619189     11.04973
US           5.2523002     23.87314
```

We can say that random forest worsens the Test MSE error to 2.837792.

"Price", "ShelveLoc" and "Age" are still most important predictor variables.

Problem 3 (Chapter 8 Exercises 9):

This problem involves the OJ data set which is part of the ISLR package.

```
> library(ISLR)
> attach(OJ)
> set.seed(1)
> summary(OJ)
```

Purchase	WeekofPurchase	StoreID	PriceCH	PriceMM
CH:653	Min. :227.0	Min. :1.00	Min. :1.690	Min. :1.690
MM:417	1st Qu.:240.0	1st Qu.:2.00	1st Qu.:1.790	1st Qu.:1.990
	Median :257.0	Median :3.00	Median :1.860	Median :2.090
	Mean :254.4	Mean :3.96	Mean :1.867	Mean :2.085
	3rd Qu.:268.0	3rd Qu.:7.00	3rd Qu.:1.990	3rd Qu.:2.180
	Max. :278.0	Max. :7.00	Max. :2.090	Max. :2.290

DiscCH	DiscMM	SpecialCH	SpecialMM
Min. :0.00000	Min. :0.00000	Min. :0.00000	Min. :0.00000
1st Qu.:0.00000	1st Qu.:0.00000	1st Qu.:0.00000	1st Qu.:0.00000
Median :0.00000	Median :0.00000	Median :0.00000	Median :0.00000
Mean :0.05186	Mean :0.1234	Mean :0.1477	Mean :0.1617
3rd Qu.:0.00000	3rd Qu.:0.2300	3rd Qu.:0.00000	3rd Qu.:0.00000
Max. :0.50000	Max. :0.8000	Max. :1.0000	Max. :1.0000

LoyalCH	SalePriceMM	SalePriceCH	PriceDiff	Store7
Min. :0.000011	Min. :1.190	Min. :1.390	Min. : -0.6700	No :714
1st Qu.:0.325257	1st Qu.:1.690	1st Qu.:1.750	1st Qu.: 0.0000	Yes:356
Median :0.600000	Median :2.090	Median :1.860	Median : 0.2300	
Mean :0.565782	Mean :1.962	Mean :1.816	Mean : 0.1465	
3rd Qu.:0.850873	3rd Qu.:2.130	3rd Qu.:1.890	3rd Qu.: 0.3200	
Max. :0.999947	Max. :2.290	Max. :2.090	Max. : 0.6400	

PctDiscMM	PctDiscCH	ListPriceDiff	STORE
Min. :0.0000	Min. :0.000000	Min. :0.0000	Min. :0.0000
1st Qu.:0.0000	1st Qu.:0.000000	1st Qu.:0.140	1st Qu.:0.0000
Median :0.0000	Median :0.000000	Median :0.240	Median :2.000
Mean :0.0593	Mean :0.02731	Mean :0.218	Mean :1.631
3rd Qu.:0.1127	3rd Qu.:0.000000	3rd Qu.:0.300	3rd Qu.:3.000
Max. :0.4020	Max. :0.25269	Max. :0.440	Max. :4.000

(a) Create a training set containing a random sample of 800 observations, and a test set containing the remaining observations.

```
> train = sample(dim(OJ)[1], 800)
> OJ.train = OJ[train, ]
> OJ.test = OJ[-train, ]
```

(b) Fit a tree to the training data, with Purchase as the response and the other variables except for Buy as predictors. Use the summary() function to produce summary statistics about the tree, and describe the results obtained. What is the training error rate? How many terminal nodes does the tree have?

```
> library(tree)
> oj.tree = tree(Purchase ~ ., data = OJ.train)
> summary(oj.tree)

Classification tree:
tree(formula = Purchase ~ ., data = OJ.train)
Variables actually used in tree construction:
[1] "LoyalCH"      "PriceDiff"    "SpecialCH"    "ListPriceDiff"
Number of terminal nodes:  8
Residual mean deviance:  0.7305 = 578.6 / 792
Misclassification error rate: 0.165 = 132 / 800
```

The tree uses variables are "LoyalCH", "PriceDiff", "SpecialCH", "ListPriceDiff" and has 8 terminal nodes and a misclassification (training) error rate of 0.165.

(c) Type in the name of the tree object in order to get a detailed text output. Pick one of the terminal nodes, and interpret the information displayed.

```
> oj.tree

node), split, n, deviance, yval, (yprob)
* denotes terminal node

1) root 800 1064.00 CH ( 0.61750 0.38250 )
2) LoyalCH < 0.508643 350 409.30 MM ( 0.27143 0.72857 )
4) LoyalCH < 0.264232 166 122.10 MM ( 0.12048 0.87952 )
8) LoyalCH < 0.0356415 57 10.07 MM ( 0.01754 0.98246 ) *
9) LoyalCH > 0.0356415 109 100.90 MM ( 0.17431 0.82569 ) *
5) LoyalCH > 0.264232 184 248.80 MM ( 0.40761 0.59239 )
10) PriceDiff < 0.195 83 91.66 MM ( 0.24096 0.75904 )
20) SpecialCH < 0.5 70 60.89 MM ( 0.15714 0.84286 ) *
21) SpecialCH > 0.5 13 16.05 CH ( 0.69231 0.30769 ) *
11) PriceDiff > 0.195 101 139.20 CH ( 0.54455 0.45545 ) *
3) LoyalCH > 0.508643 450 318.10 CH ( 0.88667 0.11333 )
6) LoyalCH < 0.764572 172 188.90 CH ( 0.76163 0.23837 )
12) ListPriceDiff < 0.235 70 95.61 CH ( 0.57143 0.42857 ) *
13) ListPriceDiff > 0.235 102 69.76 CH ( 0.89216 0.10784 ) *
7) LoyalCH > 0.764572 278 86.14 CH ( 0.96403 0.03597 ) *
```

The asterisk (*) in the line shows that it is a terminal node. I picked terminal node label 9. The splitting value of this node is 0.0356415. There are 109 subtree below this node with deviance of 100.9. The prediction at this node is Sales = MM. About 17.431% of the observations in this node take the value of CH and remaining 82.569% take the value of MM.

(d) Create a plot of the tree, and interpret the results.

```
> plot(oj.tree, cex = 0.2)
> text(oj.tree, pretty = 0, cex = 0.7)
> dev.copy2pdf(file = "MTH522_hw7_p3d.pdf", width = 8, height = 6, out.type = "pdf")
> dev.off()
```

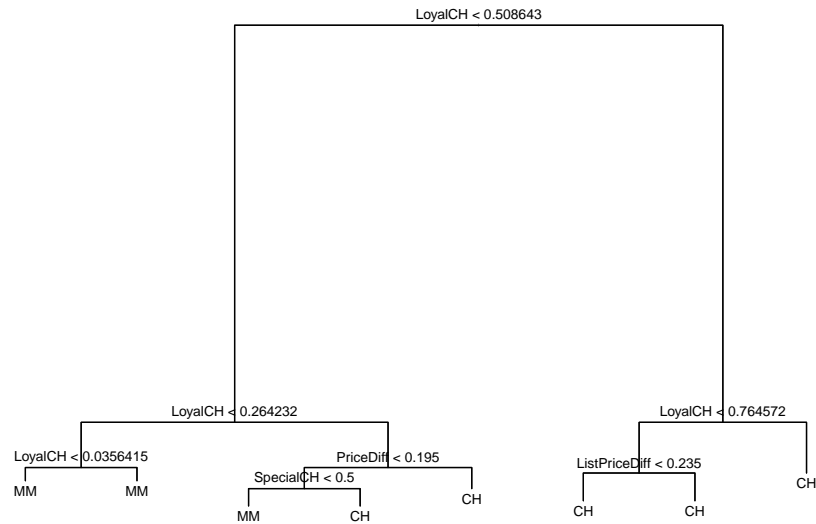


Figure 5: .

"LoyalCH" is the most important indicator of tree. The top three nodes contain "LoyalCH". If "LoyalCH" ≤ 0.264243 , the tree predicts MM. If "LoyalCH" ≥ 0.764572 , the tree predicts CH. Between those values of "LoyalCH", the decision depends on the value of "PriceDiff" and "ListPriceDiff"

(e) Predict the response on the test data, and produce a confusion matrix comparing the test labels to the predicted test labels. What is the test error rate?

```
> oj.pred = predict(oj.tree, OJ.test, type = "class")
> table(OJ.test$Purchase, oj.pred)
```

```
oj.pred
CH  MM
CH 147 12
MM  49 62
```

```
> 1 - (147 + 62) / (147+12+49+62)
[1] 0.2259259
```

Test error rate : 0.2259259

(f) Apply the `cv.tree()` function to the training set in order to determine the optimal tree size.

```
> cv.oj = cv.tree(oj.tree, FUN = prune.tree)
> cv.oj
```

```
$size
[1] 8 7 6 5 4 3 2 1
```

```
$dev
[1] 689.1001 685.8030 654.9314 653.7774 666.8890 721.2494 733.6936 1066.6499
```

```
$k
[1] -Inf 11.20965 14.72877 17.88334 23.55203 38.37537 43.02529 337.08200
```

```
$method
[1] "deviance"
```

```
attr("class")
[1] "prune" "tree.sequence"
```

(g) Produce a plot with tree size on the x-axis and cross-validated classification error rate on the y-axis.

```
> plot(cv.oj$size, cv.oj$dev, type = "b", xlab = "Tree Size", ylab = "Deviance")
> dev.copy2pdf(file = "MTH522_hw7_p3g.pdf", width = 8, height = 6, out.type = "pdf")
> dev.off()
```

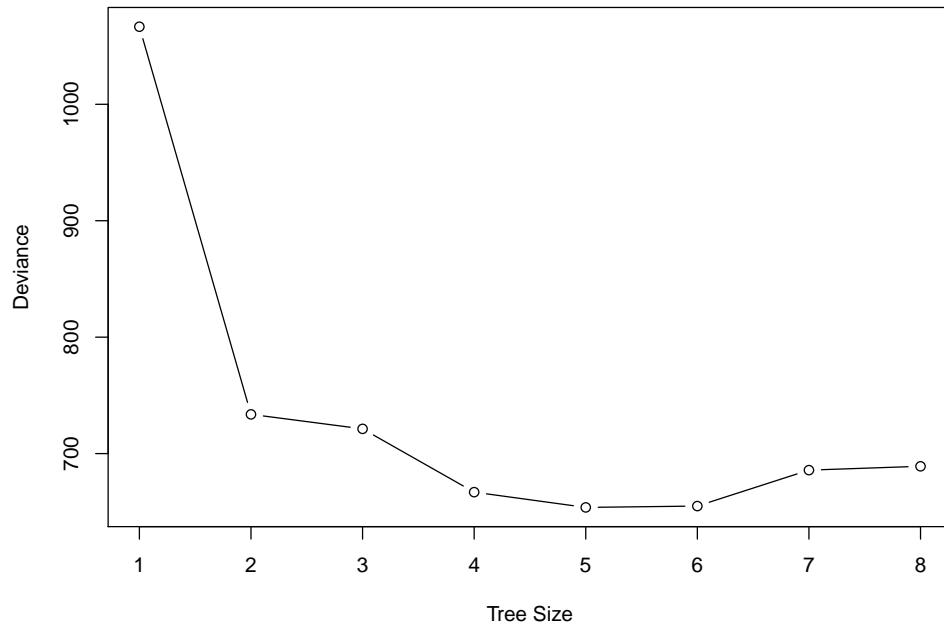


Figure 6: .

(h) Which tree size corresponds to the lowest cross-validated classification error rate?

The 5-node tree is the smallest tree with the lowest cross-validation error rate.

(i) Produce a pruned tree corresponding to the optimal tree size obtained using cross-validation. If cross-validation does not lead to selection of a pruned tree, then create a pruned tree with five terminal nodes.

```
> oj.prune = prune.tree(oj.tree, best = 5)
> plot(oj.prune)
> text(oj.prune, pretty = 0)
> dev.copy2pdf(file = "MTH522_hw7_p3i.pdf", width = 8, height = 6, out.type = "pdf")
> dev.off()
```

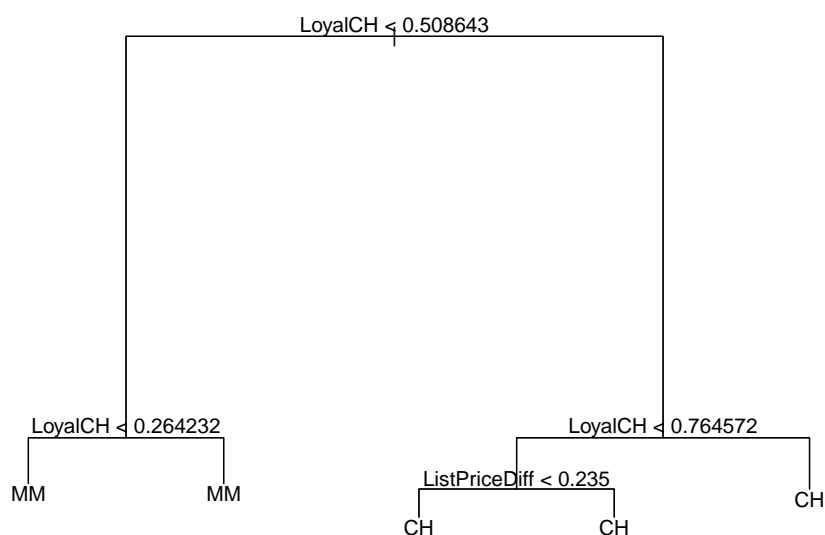


Figure 7: .

(j) Compare the training error rates between the pruned and un- pruned trees. Which is higher?

```
> summary(oj.tree)

Classification tree:
tree(formula = Purchase ~ ., data = OJ.train)
Variables actually used in tree construction:
[1] "LoyalCH"      "PriceDiff"    "SpecialCH"    "ListPriceDiff"
Number of terminal nodes:  8
Residual mean deviance:  0.7305 = 578.6 / 792
Misclassification error rate: 0.165 = 132 / 800

> summary(oj.prune)

Classification tree:
snip.tree(tree = oj.tree, nodes = 4:5)
Variables actually used in tree construction:
[1] "LoyalCH"      "ListPriceDiff"
Number of terminal nodes:  5
Residual mean deviance:  0.7829 = 622.4 / 795
Misclassification error rate: 0.1825 = 146 / 800
```

Misclassification error of prune tree (0.1825) is slightly higher than original tree (0.165).

(k) Compare the test error rates between the pruned and unpruned trees. Which is higher?

```
> pred.unprune = predict(oj.tree, OJ.test, type = "class")
> table(pred.unprune, OJ.test$Purchase)

pred.unprune  CH  MM
CH 147  49
MM  12  62

> 1 - (147+62) / (147+49+12+62)
[1] 0.2259259

> pred.prune = predict(oj.prune, OJ.test, type = "class")
> table(pred.prune, OJ.test$Purchase)

pred.prune  CH  MM
CH 119  30
MM  40  81

> 1 - (119 + 81) / (119+30+40+81)
[1] 0.2592593
```

Prune test error rate 0.2592593 is slightly higher than unpruned trees error rate of 0.2259259 .

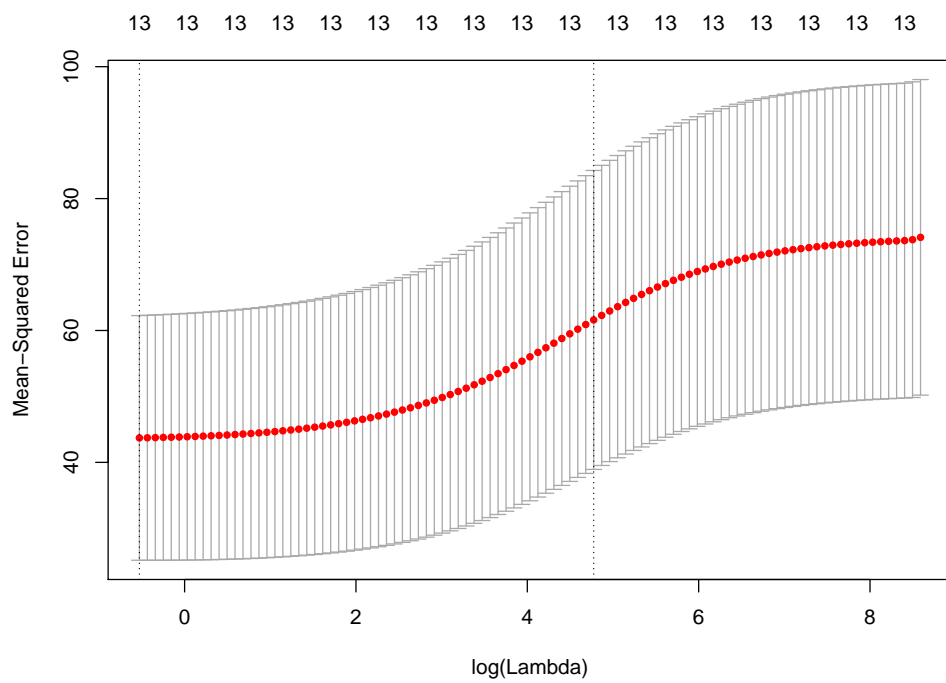


Figure 8: .

deeneme

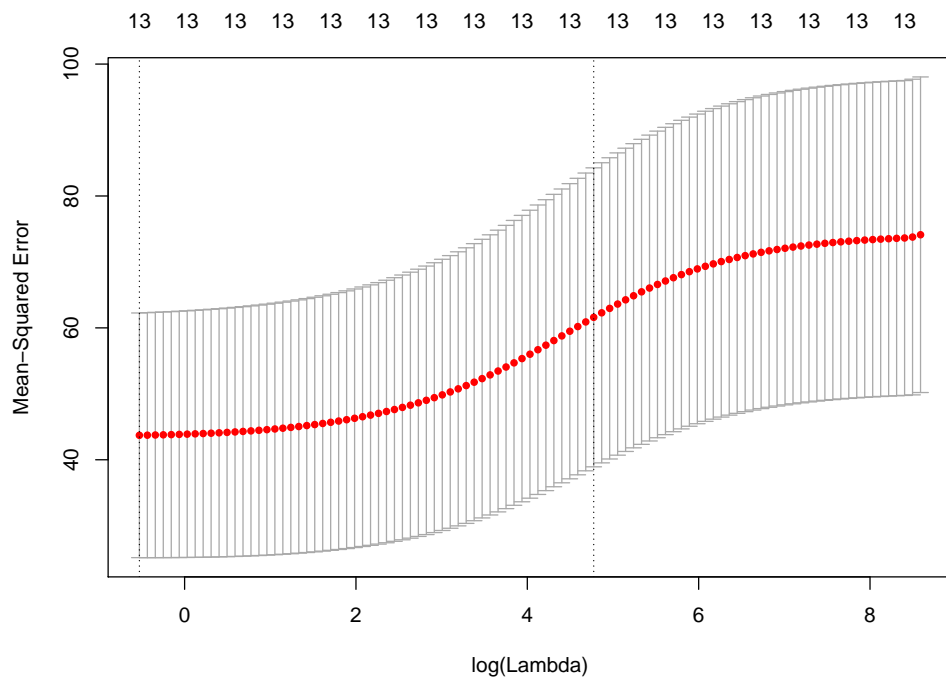


Figure 9: .