

## Solution to Homework #4—MTH 522

### Problem 1 (Chapter 5 Exercises 5):

In Chapter 4, we used logistic regression to predict the probability of default using income and balance on the Default data set. We will now estimate the test error of this logistic regression model using the validation set approach. Do not forget to set a random seed before beginning your analysis.

(a) Fit a logistic regression model that uses income and balance to predict default.

```
> library(ISLR)
> summary(Default)

default      student      balance      income
No :9667   No :7056   Min.   :  0.0   Min.   : 772
Yes: 333   Yes:2944   1st Qu.: 481.7   1st Qu.:21340
Median : 823.6   Median :34553
Mean   : 835.4   Mean   :33517
3rd Qu.:1166.3   3rd Qu.:43808
Max.   :2654.3   Max.   :73554

> attach(Default)
> set.seed(1)
> fit.glm = glm(default ~ income + balance, data = Default, family = "binomial")
> summary(fit.glm)

Call:
glm(formula = default ~ income + balance, family = "binomial",
    data = Default)

Deviance Residuals:
    Min       1Q   Median       3Q      Max
-2.4725  -0.1444  -0.0574  -0.0211   3.7245

Coefficients:
(Intercept) -1.154e+01  4.348e-01 -26.545  < 2e-16 ***
income       2.081e-05  4.985e-06   4.174  2.99e-05 ***
balance      5.647e-03  2.274e-04  24.836  < 2e-16 ***
---
Signif. codes:  0 *** 0.001 ** 0.01 * 0.05 . 0.1 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 2920.6  on 9999  degrees of freedom
Residual deviance: 1579.0  on 9997  degrees of freedom
AIC: 1585

Number of Fisher Scoring iterations: 8
```

(b) Using the validation set approach, estimate the test error of this model. In order to do this, you must perform the following steps:

(b) i. Split the sample set into a training set and a validation set.

```
> train = sample(dim(Default)[1], dim(Default)[1] / 2)
```

(b) ii. Fit a multiple logistic regression model using only the training observations.

```
> fit.glm = glm(default ~ income + balance, data = Default, family = "binomial", subset = train)
> summary(fit.glm)
```

Call:

```
glm(formula = default ~ income + balance, family = "binomial",
data = Default, subset = train)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.3583	-0.1268	-0.0475	-0.0165	3.8116

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	-1.208e+01	6.658e-01	-18.148	<2e-16 ***
income	1.858e-05	7.573e-06	2.454	0.0141 *
balance	6.053e-03	3.467e-04	17.457	<2e-16 ***

---

Signif. codes: 0 \*\*\* 0.001 \*\* 0.01 \* 0.05 . 0.1 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1457.0 on 4999 degrees of freedom

Residual deviance: 734.4 on 4997 degrees of freedom

AIC: 740.4

Number of Fisher Scoring iterations: 8

(b) iii. Obtain a prediction of default status for each individual in the validation set by computing the posterior probability of default for that individual, and classifying the individual to the default category if the posterior probability is greater than 0.5.

```
> probs = predict(fit.glm, newdata = Default[-train, ], type = "response")
> pred.glm = rep("No", length(probs))
> pred.glm[probs > 0.5] = "Yes"
```

(b) iv. Compute the validation set error, which is the fraction of the observations in the validation set that are misclassified.

```
> mean(pred.glm != Default[-train, ]$default)
[1] 0.0286
```

Using the validation set approach, the test error rate is 2.86%.

(c) Repeat the process in (b) three times, using three different splits of the observations into a training set and a validation set. Comment on the results obtained.

```
> train = sample(dim(Default)[1], dim(Default)[1] / 2)
> fit.glm = glm(default ~ income+ balance, data=Default, family="binomial", subset=train)
> probs = predict(fit.glm, newdata = Default[-train, ], type = "response")
> pred.glm = rep("No", length(probs))
> pred.glm[probs > 0.5] = "Yes"
> mean(pred.glm != Default[-train, ]$default)
[1] 0.0236

> train = sample(dim(Default)[1], dim(Default)[1] / 2)
> fit.glm = glm(default ~ income+ balance, data=Default, family="binomial", subset=train)
> probs = predict(fit.glm, newdata = Default[-train, ], type = "response")
> pred.glm = rep("No", length(probs))
> pred.glm[probs > 0.5] = "Yes"
> mean(pred.glm != Default[-train, ]$default)
[1] 0.028

> train = sample(dim(Default)[1], dim(Default)[1] / 2)
> fit.glm = glm(default ~ income+ balance, data=Default, family="binomial", subset=train)
> probs = predict(fit.glm, newdata = Default[-train, ], type = "response")
> pred.glm = rep("No", length(probs))
> pred.glm[probs > 0.5] = "Yes"
> mean(pred.glm != Default[-train, ]$default)
[1] 0.0268
```

Each time we get different test error rate. The average test error rate is 2.613333%.

(d) Now consider a logistic regression model that predicts the probability of default using income, balance, and a dummy variable for student. Estimate the test error for this model using the validation set approach. Comment on whether or not including a dummy variable for student leads to a reduction in the test error rate.

using :income, balance, student

```
> train = sample(dim(Default)[1], dim(Default)[1] / 2)
> fit.glm=glm(default~income+balance+student,data=Default,family="binomial",subset=train)
> probs = predict(fit.glm, newdata = Default[-train, ], type = "response")
> pred.glm = rep("No", length(probs))
> pred.glm[probs > 0.5] = "Yes"
> mean(pred.glm != Default[-train, ]$default)

[1] 0.0264
```

The test error rate is 2.64%, including a dummy variable for student didn't leads to a reduction in the test error rate.

## Problem 2 (Chapter 5 Exercises 6):

We continue to consider the use of a logistic regression model to predict the probability of default using income and balance on the Default data set. In particular, we will now compute estimates for the standard errors of the income and balance logistic regression coefficients in two different ways:

]

(1) using the bootstrap, and

(2) using the standard formula for computing the standard errors in the `glm()` function.

Do not forget to set a random seed before beginning your analysis.

(a) Using the `summary()` and `glm()` functions, determine the estimated standard errors for the coefficients associated with income and balance in a multiple logistic regression model that uses both predictors.

```
> set.seed(1)
> attach(Default)
The following objects are masked from Default (pos = 3):
balance, default, income, student

> fit.glm = glm(default ~ income + balance, data = Default, family = "binomial")
> summary(fit.glm)

Call:
glm(formula = default ~ income + balance, family = "binomial",
    data = Default)

Deviance Residuals:
    Min       1Q   Median       3Q      Max
-2.472542783  -0.144435943  -0.057366212  -0.021063920   3.724454436

Coefficients:
              Estimate      Std. Error  z value  Pr(>|z|)
(Intercept) -1.15404684e+01  4.34756357e-01 -26.54468 < 2.22e-16 ***
income       2.08089755e-05  4.98516718e-06   4.17418 2.9906e-05 ***
balance      5.64710294e-03  2.27373142e-04  24.83628 < 2.22e-16 ***
---
Signif. codes:  0 *** 0.001 ** 0.01 * 0.05 . 0.1 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 2920.649711  on 9999  degrees of freedom
Residual deviance: 1578.966270  on 9997  degrees of freedom
AIC: 1584.96627

Number of Fisher Scoring iterations: 8
```

The `glm()` functions estimated standard errors for the coefficients are 4.34756357e-01 , 4.98516718e-06 and 2.27373142e-04

(b) Write a function, `boot.fn()`, that takes as input the Default data set as well as an index of the observations, and that outputs the coefficient estimates for income and balance in the multiple logistic regression model.

```
> boot.fn = function(data, index) {
+ fit = glm(default ~ income + balance, data = data, family = "binomial", subset = index)
+ return (coef(fit))
+ }
```

(c) Use the `boot()` function together with your `boot.fn()` function to estimate the standard errors of the logistic regression coefficients for income and balance.

```
> library(boot)
> boot(Default, boot.fn, 300)

ORDINARY NONPARAMETRIC BOOTSTRAP

Call:
boot(data = Default, statistic = boot.fn, R = 300)

Bootstrap Statistics :
      original      bias      std. error
t1* -1.154047e+01 -6.310201e-02 4.424387e-01
t2*  2.080898e-05 -1.442808e-07 4.886839e-06
t3*  5.647103e-03  3.762129e-05 2.301732e-04
```

The `boot.fn()` function estimates the standard errors of the logistic regression coefficients are 4.424387e-01, 4.886839e-06 and 2.301732e-04

d) Comment on the estimated standard errors obtained using the `glm()` function and using your bootstrap function.

The estimated standard errors obtained using the `glm()` function and using the bootstrap function are pretty close.

$\beta_0$	$\beta_1$		$\beta_2$
4.34756357e-01	4.98516718e-06	and	2.27373142e-04
4.424387e-01	4.886839e-06	and	2.301732e-04

**Problem 3** (Chapter 5 Exercises 7):

In Sections 5.3.2 and 5.3.3, we saw that the `cv.glm()` function can be used in order to compute the LOOCV test error estimate. Alternatively, one could compute those quantities using just the `glm()` and `predict.glm()` functions, and a for loop. You will now take this approach in order to compute the LOOCV error for a simple logistic regression model on the Weekly data set. Recall that in the context of classification problems, the LOOCV error is given in (5.4).

```
> library(ISLR)
> set.seed(1)
> attach(Weekly)
> summary(Weekly)
```

Year	Lag1	Lag2	Lag3
Min. :1990	Min. : -18.1950	Min. : -18.1950	Min. : -18.1950
1st Qu.:1995	1st Qu.: -1.1540	1st Qu.: -1.1540	1st Qu.: -1.1580
Median :2000	Median : 0.2410	Median : 0.2410	Median : 0.2410
Mean :2000	Mean : 0.1506	Mean : 0.1511	Mean : 0.1472
3rd Qu.:2005	3rd Qu.: 1.4050	3rd Qu.: 1.4090	3rd Qu.: 1.4090
Max. :2010	Max. : 12.0260	Max. : 12.0260	Max. : 12.0260

  

Lag4	Lag5	Volume	Today	Direction
Min. : -18.1950	Min. : -18.1950	Min. :0.08747	Min. : -18.1950	Down:484
1st Qu.: -1.1580	1st Qu.: -1.1660	1st Qu.:0.33202	1st Qu.: -1.1540	Up :605
Median : 0.2380	Median : 0.2340	Median :1.00268	Median : 0.2410	
Mean : 0.1458	Mean : 0.1399	Mean :1.57462	Mean : 0.1499	
3rd Qu.: 1.4090	3rd Qu.: 1.4050	3rd Qu.:2.05373	3rd Qu.: 1.4050	
Max. : 12.0260	Max. : 12.0260	Max. :9.32821	Max. : 12.0260	

(a) Fit a logistic regression model that predicts Direction using Lag1 and Lag2.

```
> fit.glm = glm(Direction ~ Lag1 + Lag2, data = Weekly, family = "binomial")
> summary(fit.glm)

Call:
glm(formula = Direction ~ Lag1 + Lag2, family = "binomial", data = Weekly)

Deviance Residuals:
Min       1Q   Median       3Q      Max
-1.623  -1.261   1.001   1.083   1.506

Coefficients:
(Intercept)  0.22122    0.06147    3.599 0.000319 ***
Lag1         -0.03872    0.02622   -1.477 0.139672
Lag2          0.06025    0.02655    2.270 0.023232 *
---
Signif. codes:  0 *** 0.001 ** 0.01 * 0.05 . 0.1 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1496.2  on 1088  degrees of freedom
Residual deviance: 1488.2  on 1086  degrees of freedom
AIC: 1494.2

Number of Fisher Scoring iterations: 4
```



(b) Fit a logistic regression model that predicts Direction using Lag1 and Lag2 using all but the first observation.

```
> fit.glm.b = glm(Direction ~ Lag1 + Lag2, data = Weekly[-1, ], family = "binomial")
> summary(fit.glm.b)

Call:
glm(formula = Direction ~ Lag1 + Lag2, family = "binomial", data = Weekly[-1,
])

Deviance Residuals:
Min       1Q   Median       3Q      Max
-1.6258  -1.2617   0.9999   1.0819   1.5071

Coefficients:
(Intercept)  0.22324    0.06150    3.630 0.000283 ***
Lag1         -0.03843    0.02622   -1.466 0.142683
Lag2          0.06085    0.02656    2.291 0.021971 *
---
Signif. codes:  0 *** 0.001 ** 0.01 * 0.05 . 0.1 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1494.6  on 1087  degrees of freedom
Residual deviance: 1486.5  on 1085  degrees of freedom
AIC: 1492.5

Number of Fisher Scoring iterations: 4
```

(c) Use the model from (b) to predict the direction of the first observation. You can do this by predicting that the first observation will go up if  $P(\text{Direction} = \text{"Up"} | \text{Lag1, Lag2}) \geq 0.5$ . Was this observation correctly classified?

```
> predict.glm(fit.glm.b, Weekly[1, ], type = "response") > 0.5
1
TRUE
```

Prediction for the observation is Up. This observation was not correctly classified; true direction is Down.

(d) Write a for loop from  $i=1$  to  $i=n$ , where  $n$  is the number of observations in the data set, that performs each of the following steps:

- i. Fit a logistic regression model using all but the  $i$ th observation to predict Direction using Lag1 and Lag2.
- ii. Compute the posterior probability of the market moving up for the  $i$ th observation.
- iii. Use the posterior probability for the  $i$ th observation in order to predict whether or not the market moves up.
- iv. Determine whether or not an error was made in predicting the direction for the  $i$ th observation. If an error was made, then indicate this as a 1, and otherwise indicate it as a 0.

```
> error = rep(0, dim(Weekly)[1])
> for (i in 1:dim(Weekly)[1]) {
+   fit.glm = glm(Direction ~ Lag1 + Lag2, data = Weekly[-i, ], family = "binomial")
+   pred.up = predict.glm(fit.glm, Weekly[i, ], type = "response") > 0.5
+   true.up = Weekly[i, ]$Direction == "Up"
+   if (pred.up != true.up)
+     error[i] = 1
+ }
> error
[1] 1 1 0 1 0 1 0 0 0 1 1 0 1 0 1 0 1 0 1 0 0 0 1 1 1 1 1 1 0 1 1 1 1 0 1 0 0
[38] 0 1 0 1 0 0 1 0 1 1 1 0 1 0 0 0 1 0 0 1 1 0 0 0 0 1 0 1 1 0 0 1 0 1 1 0 0
[75] 0 1 0 1 1 0 0 1 1 0 1 1 0 0 1 0 0 1 1 1 0 0 0 0 0 1 0 1 1 0 0 1 0 1 0 0 1
[112] 1 0 0 1 0 0 1 0 0 1 1 1 1 0 0 0 1 0 1 0 1 1 0 0 0 1 1 1 0 0 0 1 0 0 0 0 0
[149] 0 1 1 1 0 1 0 0 1 1 0 1 0 0 1 1 0 0 1 0 0 1 0 0 1 1 1 0 1 0 1 0 0 0 0 0 0
[186] 0 0 1 1 0 1 0 1 0 1 0 1 0 0 1 0 0 1 0 0 1 0 0 1 0 1 0 1 1 1 0 0 1 1 0 1 0 0 1 1
[223] 0 0 0 1 1 1 0 1 0 1 0 1 0 0 0 1 1 0 1 0 1 0 1 0 1 0 1 1 0 1 0 1 0 0 1 0 0 1 0
[260] 0 0 0 0 1 0 0 0 1 0 0 1 0 0 0 1 0 0 1 0 0 1 0 0 1 0 0 1 0 1 1 0 0 0 0 0 1 0 1 0
[297] 0 1 0 0 0 1 0 0 1 1 0 0 1 0 0 0 0 1 0 1 1 0 0 1 0 1 0 1 1 0 0 0 1 0 1 0 0
[334] 1 1 1 1 0 1 0 0 1 0 0 0 1 0 1 0 1 0 0 0 0 0 1 1 0 0 1 0 0 1 0 0 1 0 0 0 1 1 0 1
[371] 1 1 1 1 0 0 0 1 0 0 0 0 0 0 0 1 0 1 1 0 0 1 1 0 0 0 0 0 1 0 0 1 1 1 0 1 0 1
[408] 0 1 1 1 0 1 0 0 0 0 0 0 0 0 0 0 1 0 1 0 1 0 1 0 0 1 0 1 0 0 0 0 0 1 1 1 1
[445] 0 1 1 0 1 0 1 1 0 1 0 0 1 0 0 1 1 0 0 0 0 1 1 0 0 1 0 1 0 0 0 1 0 0 1 0 0
[482] 0 1 1 1 0 1 0 0 0 1 0 1 1 1 0 0 0 0 1 1 1 0 1 1 0 1 0 1 0 0 0 1 0 1 0 0 0 1 0
[519] 1 1 0 0 1 1 0 0 0 1 1 0 1 0 1 1 1 1 1 0 0 0 1 0 0 0 1 1 0 1 0 0 0 1 1 1 1
[556] 1 1 0 1 0 1 0 0 1 0 0 1 1 1 0 0 0 1 1 1 1 1 1 1 1 1 1 1 0 1 0 0 1 0 0 1 0 1
[593] 0 0 1 0 0 1 0 1 1 0 1 1 1 0 1 0 1 0 1 0 0 0 1 0 1 0 1 0 1 0 1 1 0 1 1 0 1 0 1
[630] 0 1 1 1 1 0 1 1 0 0 0 1 1 1 1 0 1 1 1 0 1 0 0 0 1 1 1 1 1 1 0 1 0 0 1 0 0
[667] 0 1 1 0 1 0 1 1 1 1 0 0 0 1 1 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 1 0 0 1 0 1 1
[704] 0 0 0 0 1 0 1 0 1 0 1 0 0 1 1 0 0 0 0 0 0 0 0 0 0 0 1 0 0 1 0 0 1 1 1 0 1 1 0
[741] 1 1 1 1 0 0 0 1 1 1 1 1 1 0 1 0 0 0 0 0 0 1 0 1 1 1 0 0 0 1 0 0 1 0 0 0 1
[778] 1 1 0 0 0 1 0 0 1 1 1 0 0 1 0 1 0 1 0 0 1 0 0 1 0 0 0 0 0 1 0 1 1 0 0 1 1
[815] 0 1 1 1 0 0 0 0 0 1 1 0 0 1 0 0 1 0 1 0 0 0 1 1 0 1 1 0 1 0 1 0 1 1 0 0 1
[852] 1 1 0 1 1 0 0 0 1 0 1 0 1 0 1 0 0 0 0 0 1 0 0 1 1 0 0 1 0 1 0 1 1 0 1 0 1
[889] 1 0 0 0 1 0 0 0 0 0 0 0 0 1 0 1 0 1 0 0 0 1 1 1 1 0 1 1 0 0 0 0 0 1 0 0 1
[926] 0 0 0 0 1 0 1 1 1 0 0 1 1 0 1 1 1 1 0 1 0 1 0 1 0 1 0 1 0 1 0 0 1 1 1 1 0 1
[963] 0 0 0 1 1 1 0 1 1 1 1 0 0 0 0 1 1 0 0 0 0 1 0 0 1 1 1 0 0 1 1 1 0 1 0 0 0
[1000] 0 1 0 0 1 0 1 0 0 1 1 1 1 0 1 0 0 1 0 0 1 0 0 1 1 0 1 1 1 0 1 1 0 0 0 1 0
[1037] 1 0 1 1 1 1 0 0 1 0 0 0 0 0 0 1 0 1 1 0 0 0 0 0 1 1 1 0 0 0 1 0 1 1 1 0 0
[1074] 0 0 1 0 0 0 0 0 1 0 1 0 0 0 0 0
```

```
> sum(error)
[1] 490
```

(e) Take the average of the  $n$  numbers obtained in (d)iv in order to obtain the LOOCV estimate for the test error. Comment on the results.

```
> mean(error)
[1] 0.4499541
```

The LOOCV estimate for the test error rate is 44.99541%.

**Problem 4** (Chapter 5 Exercises 9 ):

We will now consider the Boston housing data set, from the MASS library.

```
> library(MASS)
> set.seed(1)
> attach(Boston)
> summary(Boston)
```

crim	zn	indus	chas
Min. : 0.00632	Min. : 0.00	Min. : 0.46	Min. : 0.00000
1st Qu.: 0.08204	1st Qu.: 0.00	1st Qu.: 5.19	1st Qu.: 0.00000
Median : 0.25651	Median : 0.00	Median : 9.69	Median : 0.00000
Mean : 3.61352	Mean : 11.36	Mean : 11.14	Mean : 0.06917
3rd Qu.: 3.67708	3rd Qu.: 12.50	3rd Qu.: 18.10	3rd Qu.: 0.00000
Max. : 88.97620	Max. : 100.00	Max. : 27.74	Max. : 1.00000

  

nox	rm	age	dis
Min. : 0.3850	Min. : 3.561	Min. : 2.90	Min. : 1.130
1st Qu.: 0.4490	1st Qu.: 5.886	1st Qu.: 45.02	1st Qu.: 2.100
Median : 0.5380	Median : 6.208	Median : 77.50	Median : 3.207
Mean : 0.5547	Mean : 6.285	Mean : 68.57	Mean : 3.795
3rd Qu.: 0.6240	3rd Qu.: 6.623	3rd Qu.: 94.08	3rd Qu.: 5.188
Max. : 0.8710	Max. : 8.780	Max. : 100.00	Max. : 12.127

  

rad	tax	ptratio	black
Min. : 1.000	Min. : 187.0	Min. : 12.60	Min. : 0.32
1st Qu.: 4.000	1st Qu.: 279.0	1st Qu.: 17.40	1st Qu.: 375.38
Median : 5.000	Median : 330.0	Median : 19.05	Median : 391.44
Mean : 9.549	Mean : 408.2	Mean : 18.46	Mean : 356.67
3rd Qu.: 24.000	3rd Qu.: 666.0	3rd Qu.: 20.20	3rd Qu.: 396.23
Max. : 24.000	Max. : 711.0	Max. : 22.00	Max. : 396.90

  

lstat	medv
Min. : 1.73	Min. : 5.00
1st Qu.: 6.95	1st Qu.: 17.02
Median : 11.36	Median : 21.20
Mean : 12.65	Mean : 22.53
3rd Qu.: 16.95	3rd Qu.: 25.00
Max. : 37.97	Max. : 50.00

(a) Based on this data set, provide an estimate for the population mean of medv. Call this estimate  $\mu$ .

```
> estimate.mu = mean(medv)
> estimate.mu

[1] 22.53281
```

(b) Provide an estimate of the standard error of  $\mu$ . Interpret this result.

Hint: We can compute the standard error of the sample mean by dividing the sample standard deviation by the square root of the number of observations.

```
> estimate.mu <- sd(medv) / sqrt(dim(Boston)[1])
> estimate.mu

[1] 0.4088611
```

(c) Now estimate the standard error of  $\mu$  using the bootstrap. How does this compare to your answer from (b)?

```
> boot.fn = function(data, index) {
+   mu = mean(data[index])
+   return (mu)
+ }
> boot(medv, boot.fn, 1000)

ORDINARY NONPARAMETRIC BOOTSTRAP

Call:
boot(data = medv, statistic = boot.fn, R = 1000)

Bootstrap Statistics :
original    bias    std. error
t1* 22.53281 0.0027583   0.4120131
```

Estimate the standard error of  $\mu$  using the bootstrap is 0.4120131 and from (b) estimate of the standard error of  $\mu$  is 0.4088611. They are similar.

(d) Based on your bootstrap estimate from (c), provide a 95 % confidence interval for the mean of medv. Compare it to the results obtained using `t.test(Boston$medv)`.

```
> t.test(medv)

One Sample t-test

data:  medv
t = 55.111, df = 505, p-value < 2.2e-16
alternative hypothesis: true mean is not equal to 0
95 percent confidence interval:
 21.72953 23.33608
sample estimates:
mean of x
 22.53281

> CI.estimate.mu = c(22.53 - 2 * 0.4119, 22.53 + 2 * 0.4119)
> CI.estimate.mu

[1] 21.7062 23.3538
```

```
The bootstrap confidence interval estimate : 21.72953 23.33608
t.test estimate : 21.7062 23.3538
Two estimates are only 0.02 away from each other
```

(e) Based on this data set, provide an estimate,  $\mu_{med}$ , for the median value of medv in the population.

```
> med.estimate <- median(medv)
> med.estimate

[1] 21.2
```

(f) We now would like to estimate the standard error of  $\mu_{med}$ . Unfortunately, there is no simple formula for computing the standard error of the median. Instead, estimate the standard error of the median using the bootstrap. Comment on your findings.

```
> boot.fn = function(data, index) {
+   mu = median(data[index])
+   return (mu)
+ }
> boot(medv, boot.fn, 1000)

ORDINARY NONPARAMETRIC BOOTSTRAP

Call:
boot(data = medv, statistic = boot.fn, R = 1000)

Bootstrap Statistics :
      original    bias      std. error
t1*         21.2 -0.0119   0.3856515
```

Estimate median value is 21.2. It is same value from (e). The standard error 0.3856515 is too small compared to median value.

(g) Based on this data set, provide an estimate for the tenth percentile of medv in Boston suburbs. Call this quantity  $\mu_{0.1}$ . (You can use the quantile() function.)

```
> percent.estimate <- quantile(medv, c(0.1))
> percent.estimate

10%
12.75
```

(h) Use the bootstrap to estimate the standard error of  $\mu_{0.1}$ . Comment on your findings.

```
> boot.fn = function(data, index) {
+   mu = quantile(data[index], c(0.1))
+   return (mu)
+ }
> boot(medv, boot.fn, 1000)

ORDINARY NONPARAMETRIC BOOTSTRAP

Call:
boot(data = medv, statistic = boot.fn, R = 1000)

Bootstrap Statistics :
      original    bias      std. error
t1*         12.75  0.0041   0.5098364
```

Estimate tenth percentile value is 12.75. It is same value from (g) estimate for the tenth percentile of medv. The standard error 0.5098364 is too small compared to percentile value.