

Solution to Homework #6—MTH 522

Problem 1 (Chapter 7 Exercises 6):

In this exercise, you will further analyze the Wage data set considered throughout this chapter. (a) Perform polynomial regression to predict "wage" using "age". Use cross-validation to select the optimal degree d for the polynomial. What degree was chosen, and how does this compare to the results of hypothesis testing using ANOVA? Make a plot of the resulting polynomial fit to the data.

Let's see the Wage dataset Plot the "wage" an "age"

```
> library(ISLR)
> library(boot)
> set.seed(1)
> summary(Wage)
```

year	age	sex	maritl	race
Min. :2003	Min. :18.00	1. Male :3000	1. Never Married: 648	1. White:2480
1st Qu.:2004	1st Qu.:33.75	2. Female: 0	2. Married :2074	2. Black: 293
Median :2006	Median :42.00		3. Widowed : 19	3. Asian: 190
Mean :2006	Mean :42.41		4. Divorced : 204	4. Other: 37
3rd Qu.:2008	3rd Qu.:51.00		5. Separated : 55	
Max. :2009	Max. :80.00			

education	region	jobclass
1. < HS Grad :268	2. Middle Atlantic :3000	1. Industrial :1544
2. HS Grad :971	1. New England : 0	2. Information:1456
3. Some College :650	3. East North Central: 0	
4. College Grad :685	4. West North Central: 0	
5. Advanced Degree:426	5. South Atlantic : 0	
6. East South Central: 0		
(Other) : 0		

health	health_ins	logwage	wage
1. <=Good : 858	1. Yes:2083	Min. :3.000	Min. : 20.09
2. >=Very Good:2142	2. No : 917	1st Qu.:4.447	1st Qu.: 85.38
Median :4.653	Median :104.92		
Mean :4.654	Mean :111.70		
3rd Qu.:4.857	3rd Qu.:128.68		
Max. :5.763	Max. :318.34		


```
> plot(age, wage)
> dev.copy2pdf(file = "MTH522_hw6_p1a_1.pdf", width = 8, height = 6, out.type = "pdf")
> dev.off()
```




Figure 1: .

Now, I will perform polynomial regression:

```
> deltas = rep(0, 10)
> for (i in 1:10) {
+   glm.fit = glm(wage~poly(age, i), data=Wage)
+   deltas[i] = cv.glm(Wage, glm.fit, K=10)$delta[2]
+ }
> plot(1:10, deltas, xlab="Degree", ylab="CV Prediction Error Estimate", type="l",
+ pch=20 )
> points(which.min(deltas), deltas[which.min(deltas)], col = "red", cex = 2, pch =20)
> dev.copy2pdf(file = "MTH522_hw6_p1a_2.pdf", width = 8, height = 6, out.type = "pdf")
> dev.off()

> which.min( deltas )

[1] 4
```

From the above plot, we can see that $d=4$ is the optimal degree for the polynomial, where curve stops decreasing and starts increasing.




Figure 2: .

We now use ANOVA approach. The `anova()` function, which performs an analysis of variance (ANOVA, using an F-test) in order to test the null hypothesis that a model M_1 is sufficient to explain the data against the alternative hypothesis that a more complex model M_2 is required (Page:290)

```

> fit.1 = lm(wage ~ age, data=Wage)
> fit.2 = lm(wage ~ poly(age, 2), data=Wage)
> fit.3 = lm(wage ~ poly(age, 3), data=Wage)
> fit.4 = lm(wage ~ poly(age, 4), data=Wage)
> fit.5 = lm(wage ~ poly(age, 5), data=Wage)
> anova(fit.1, fit.2, fit.3, fit.4, fit.5 )
Analysis of Variance Table

Model 1: wage ~ age
Model 2: wage ~ poly(age, 2)
Model 3: wage ~ poly(age, 3)
Model 4: wage ~ poly(age, 4)
Model 5: wage ~ poly(age, 5)
Res.Df    RSS Df Sum of Sq      F    Pr(>F)
1    2998 5022216
2    2997 4793430  1    228786 143.5931 < 2.2e-16 ***
3    2996 4777674  1     15756  9.8888 0.001679 **
4    2995 4771604  1      6070  3.8098 0.051046 .
5    2994 4770322  1      1283  0.8050 0.369682
---
Signif. codes:  0 *** 0.001 ** 0.01 * 0.05 . 0.1 1

```

As we learn from the text book (page:290), the p-value comparing the linear Model 1 to the quadratic Model 2 is essentially zero (0.001679), indicating that a linear fit is not sufficient. Similarly the p-value comparing the quadratic Model 2 to the cubic Model 3 is very low (0.0017), so the quadratic fit is also insufficient. The p-value comparing the cubic and degree-4 polynomials, Model 3 and Model 4, is approximately 5 % while the degree-5 polynomial Model 5 seems unnecessary because its p-value is 0.37. Hence, either a cubic or a quartic polynomial appear to provide a reasonable fit to the data, but lower or higher order models are not justified.

The polynomial prediction on the data

```
> plot(wage~age, data=Wage, col="darkgrey")
> agelims = range(Wage$age)
> age.grid = seq(from=agelims[1], to=agelims[2])
> fit = lm(wage~poly(age, 3), data=Wage)
> pred = predict(fit, data.frame(age=age.grid))
> lines(age.grid, pred, col="blue", lwd=2)
> dev.copy2pdf(file = "MTH522_hw6_p1a_1.pdf", width = 8, height = 6, out.type = "pdf")
dev.off()
> dev.copy2pdf(file = "MTH522_hw6_p1a_1.pdf", width = 8, height = 6, out.type = "pdf")
dev.off()
> plot(wage~age, data=Wage, col="darkgrey")
> agelims = range(Wage$age)
> age.grid = seq(from=agelims[1], to=agelims[2])
> fit = lm(wage~poly(age, 3), data=Wage)
> pred = predict(fit, data.frame(age=age.grid))
> lines(age.grid, pred, col="blue", lwd=2)
> dev.copy2pdf(file = "MTH522_hw6_p1a_3.pdf", width = 8, height = 6, out.type = "pdf")
> dev.off()
```




Figure 3: .

(b) Fit a step function to predict wage using age, and perform cross-validation to choose the optimal number of cuts. Make a plot of the fit obtained.

To predict wage using age, I will use cross-validation cut points of up to 10.

```
> cvs = rep(NA, 10)
> for (i in 2:10) {
+   Wage$age.cut = cut(Wage$age, i)
+   fit = glm(wage~age.cut, data=Wage)
+   cvs[i] = cv.glm(Wage, fit, K=10)$delta[2]
+ }
> plot(2:10, cvs[-1], xlab="Cuts", ylab="CV error",
+ type="l", pch=20, lwd=2)
> points(which.min(cvs), cvs[which.min(cvs)], col="blue", cex=2, pch=20)
> cvs = rep(NA, 10)
> for (i in 2:10) {
+   Wage$age.cut = cut(Wage$age, i)
+   fit = glm(wage~age.cut, data=Wage)
+   cvs[i] = cv.glm(Wage, fit, K=10)$delta[2]
+ }
> plot(2:10, cvs[-1], xlab="Cuts", ylab="CV error",
+ type="l", pch=20, lwd=2)
> points(which.min(cvs), cvs[which.min(cvs)], col="blue", cex=2, pch=20)
> dev.copy2pdf(file = "MTH522_hw6_p1b_1.pdf", width = 8, height = 6, out.type = "pdf")
> dev.off()
```




Figure 4: .

From above plot, we can see that CV error is minimum at 8 cuts. Now we can fit dataset with step function using this 8 cuts.

```
> fit = glm(wage~cut(age, 8), data=Wage)
> agelims = range(Wage$age)
> age.grid = seq(from=agelims[1], to=agelims[2])
> pred = predict(fit, data.frame(age=age.grid))
> plot(wage~age, data=Wage, col="darkgrey")
> lines(age.grid, pred, col="red", lwd=2)
> dev.copy2pdf(file = "MTH522_hw6_p1b_2.pdf", width = 8, height = 6, out.type = "pdf")
> dev.off()
```




Figure 5: .

Problem 2 (Chapter 7 Exercises 7):

The Wage data set contains a number of other features not explored in this chapter, such as marital status (maritl), job class (jobclass), and others. Explore the relationships between some of these other predictors and wage, and use non-linear fitting techniques in order to fit flexible models to the data. Create plots of the results obtained, and write a summary of your findings.

```
> library(ISLR)
> set.seed(1)
> summary(Wage)
```

year	age	sex	maritl	race
Min. :2003	Min. :18.00	1. Male :3000	1. Never Married: 648	1. White:2480
1st Qu.:2004	1st Qu.:33.75	2. Female: 0	2. Married :2074	2. Black: 293
Median :2006	Median :42.00		3. Widowed : 19	3. Asian: 190
Mean :2006	Mean :42.41		4. Divorced : 204	4. Other: 37
3rd Qu.:2008	3rd Qu.:51.00		5. Separated : 55	
Max. :2009	Max. :80.00			

education	region	jobclass
1. < HS Grad :268	2. Middle Atlantic :3000	1. Industrial :1544
2. HS Grad :971	1. New England : 0	2. Information:1456
3. Some College :650	3. East North Central: 0	
4. College Grad :685	4. West North Central: 0	
5. Advanced Degree:426	5. South Atlantic : 0	
6. East South Central: 0		
(Other) : 0		

health	health_ins	logwage	wage	age.cut
1. <=Good : 858	1. Yes:2083	Min. :3.000	Min. : 20.09	(42.8,49] :640
2. >=Very Good:2142	2. No : 917	1st Qu.:4.447	1st Qu.: 85.38	(36.6,42.8]:542
Median :4.653	Median :104.92	(30.4,36.6]:445		
Mean :4.654	Mean :111.70	(49,55.2] :441		
3rd Qu.:4.857	3rd Qu.:128.68	(24.2,30.4]:347		
Max. :5.763	Max. :318.34	(55.2,61.4]:270		
(Other) :315				


```

> par(mfrow = c(1, 2), las=2, cex.axis = 0.6)
> plot(Wage$maritl, Wage$wage)
> plot(Wage$jobclass, Wage$wage)
> dev.copy2pdf(file = "MTH522_hw6_p2_1.pdf", width = 8, height = 6, out.type = "pdf")
> dev.off()

```




Figure 6: .

Conclusion: Plot 1, A married group earns more money on average. Plot 2, Informational jobs group earns more on average.

```

> fit1 <- gam(wage ~ maritl, data = Wage)
> fit2 <- gam(wage ~ jobclass, data = Wage)
> fit3 <- gam(wage ~ maritl + jobclass, data = Wage)
> fit4 <- gam(wage ~ maritl + jobclass + s(age, 5) , data = Wage)
> fit5 <- gam(wage ~ maritl + jobclass + s(age, 5) + education , data = Wage)
> anova(fit1, fit2, fit3, fit4, fit5)
Analysis of Deviance Table

Model 1: wage ~ maritl
Model 2: wage ~ jobclass
Model 3: wage ~ maritl + jobclass
Model 4: wage ~ maritl + jobclass + s(age, 5)
Model 5: wage ~ maritl + jobclass + s(age, 5) + education
Resid. Df Resid. Dev      Df Deviance  Pr(>Chi)
1      2995      4858941
2      2998      4998547 -3.0000   -139606 < 2.2e-16 ***
3      2994      4654752  4.0000    343795 < 2.2e-16 ***
4      2989      4472687  4.9997    182065 < 2.2e-16 ***
5      2985      3604891  4.0000     867796 < 2.2e-16 ***
---
Signif. codes:  0 *** 0.001 ** 0.01 * 0.05 . 0.1 1

```

As we can see "fit5 : maritl + jobclass + s(age, 5) + education" add statistically significant improvements to the model and significantly better.

```

> fit1 <- gam(wage ~ maritl, data = Wage)
> fit2 <- gam(wage ~ jobclass, data = Wage)
> fit3 <- gam(wage ~ maritl + jobclass, data = Wage)
> fit4 <- gam(wage ~ maritl + jobclass + s(age, 5) , data = Wage)
> fit5 <- gam(wage ~ maritl + jobclass + s(age, 5) + education , data = Wage)
> anova(fit1, fit2, fit3, fit4, fit5)
Analysis of Deviance Table

Model 1: wage ~ maritl
Model 2: wage ~ jobclass
Model 3: wage ~ maritl + jobclass
Model 4: wage ~ maritl + jobclass + s(age, 5)
Model 5: wage ~ maritl + jobclass + s(age, 5) + education
Resid. Df Resid. Dev      Df Deviance  Pr(>Chi)
1      2995      4858941
2      2998      4998547 -3.0000   -139606 < 2.2e-16 ***
3      2994      4654752  4.0000    343795 < 2.2e-16 ***
4      2989      4472687  4.9997    182065 < 2.2e-16 ***
5      2985      3604891  4.0000     867796 < 2.2e-16 ***
---
Signif. codes:  0 *** 0.001 ** 0.01 * 0.05 . 0.1 1

```

```

> par(mfrow = c(2, 2), las=2, cex.axis = 0.6)
> plot(fit5, se = T, col = "blue")
> dev.copy2pdf(file = "MTH522_hw6_p2_2.pdf", width = 8, height = 6, out.type = "pdf")
> dev.off()

```

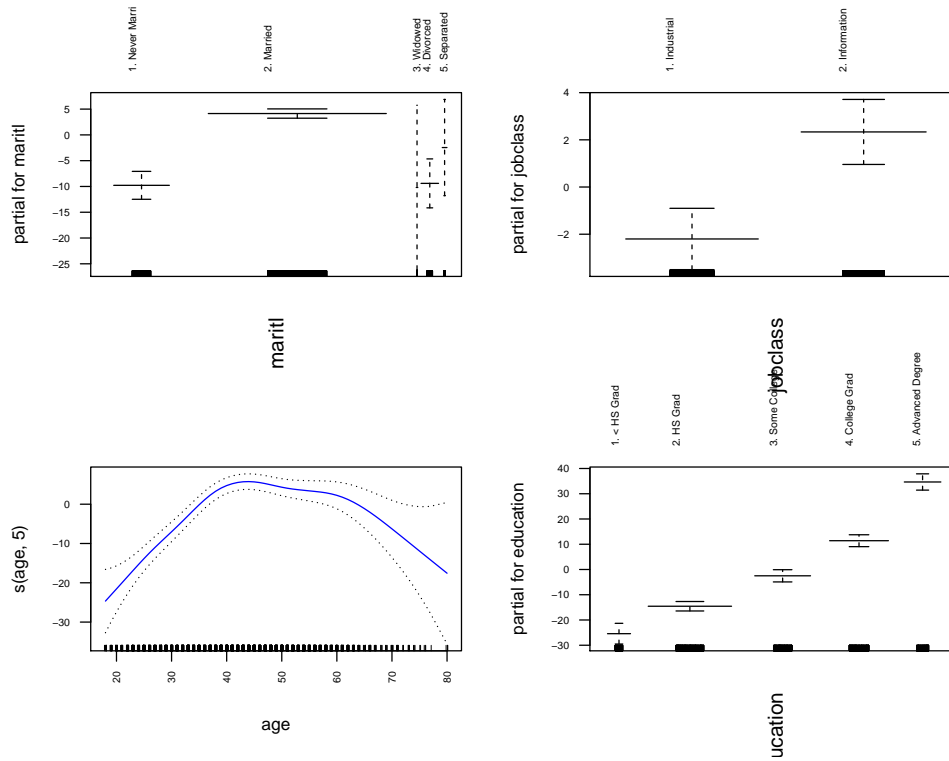


Figure 7: .

Problem 3 (Chapter 7 Exercises 9):

This question uses the variables `dis` (the weighted mean of distances to five Boston employment centers) and `nox` (nitrogen oxides concentration in parts per 10 million) from the Boston data. We will treat `dis` as the predictor and `nox` as the response.

(a) Use the `poly()` function to fit a cubic polynomial regression to predict `nox` using `dis`. Report the regression output, and plot the resulting data and polynomial fits.

```
> set.seed(1)
> library(MASS)
> attach(Boston)
> fit = lm(nox ~ poly(dis, 3), data = Boston)
> summary(fit)
```

Call:
lm(formula = nox ~ poly(dis, 3), data = Boston)

Residuals:

Min	1Q	Median	3Q	Max
-0.121130	-0.040619	-0.009738	0.023385	0.194904

Coefficients:

Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.554695	0.002759	201.021 < 2e-16 ***
poly(dis, 3)1	-2.003096	0.062071	-32.271 < 2e-16 ***
poly(dis, 3)2	0.856330	0.062071	13.796 < 2e-16 ***
poly(dis, 3)3	-0.318049	0.062071	-5.124 4.27e-07 ***

Signif. codes: 0 *** 0.001 ** 0.01 * 0.05 . 0.1 1

Residual standard error: 0.06207 on 502 degrees of freedom
Multiple R-squared: 0.7148, Adjusted R-squared: 0.7131
F-statistic: 419.3 on 3 and 502 DF, p-value: < 2.2e-16

```
> dislim = range(dis)
> grid = seq(from = dislim[1], to = dislim[2], by = 0.1)
> pred = predict(fit, list(dis = grid))
> plot(nox ~ dis, data = Boston, col = "darkgrey")
> lines(grid, pred, col = "red", lwd = 2)
> dislim = range(dis)
> grid = seq(from = dislim[1], to = dislim[2], by = 0.1)
> pred = predict(fit, list(dis = grid))
> plot(nox ~ dis, data = Boston, col = "darkgrey")
> lines(grid, pred, col = "red", lwd = 2)
> dev.copy2pdf(file = "MTH522_hw6_p3a_1.pdf", width = 8, height = 6, out.type = "pdf")
> dev.off()
```




Figure 8: .

I used the `poly()` function "`nox ~ poly(dis, 3)`" to fit a cubic polynomial regression to predict `nox` using `dis` and find out that all polynomial terms are significant while predicting `nox` using `dis`. As you can see, above figure shows that a smooth curve fitting the data well.

(b) Plot the polynomial fits for a range of different polynomial degrees (say, from 1 to 10), and report the associated residual sum of squares.

Let's plot polynomial degrees from 1 to 10 and report the associated residual sum of squares (RSS).

```
> all.rss = rep(NA, 10)
> for (i in 1:10) {
+   fit = lm(nox ~ poly(dis, i), data = Boston)
+   all.rss[i] = sum(fit$residuals^2)
+ }
> all.rss
[1] 2.768563 2.035262 1.934107 1.932981 1.915290 1.878257 1.849484 1.835630 1.833331 1.832171

> plot(1:10, all.rss, xlab = "Degree", ylab = "RSS", type = "l")
> dev.copy2pdf(file = "MTH522_hw6_p3b.pdf", width = 8, height = 6, out.type = "pdf")
> dev.off()
```




Figure 9: .

From plot, minimum for a polynomial of degree 10. RSS decreases with the degree of polynomial.

(c) Perform cross-validation or another approach to select the optimal degree for the polynomial, and explain your results.

```
> library(boot)
> all.deltas = rep(NA, 10)
> for (i in 1:10) {
+   fit = glm(nox ~ poly(dis, i), data = Boston)
+   all.deltas[i] = cv.glm(Boston, fit, K = 10)$delta[1]
+ }
> plot(1:10, all.deltas, xlab = "Degree", ylab = "CV Test error", type = "l")
> dev.copy2pdf(file = "MTH522_hw6_p3c.pdf", width = 8, height = 6, out.type = "pdf")
> dev.off()

> which.min(all.deltas)
[1] 4
```




Figure 10: .

As we may see CV test error minimizes at the optimal degree for the polynomial 4. The error reduces between 1 to 3. From 3 to 4 reduces just little bit and after 4 start increasing for higher degrees.

(d) Use the `bs()` function to fit a regression spline to predict `nox` using `dis`. Report the output for the fit using four degrees of freedom. How did you choose the knots? Plot the resulting fit.

```
> library(splines)
> fit = lm(nox ~ bs(dis, df = 4, knots = c(4, 7, 11)), data = Boston)
> summary(fit)

Call:
lm(formula = nox ~ bs(dis, df = 4, knots = c(4, 7, 11)), data = Boston)

Residuals:
Min       1Q   Median       3Q      Max
-0.124567 -0.040355 -0.008702  0.024740  0.192920

Coefficients:
Estimate Std. Error t value Pr(>|t|)
(Intercept)          0.73926    0.01331  55.537 < 2e-16 ***
bs(dis, df = 4, knots = c(4, 7, 11))1 -0.08861    0.02504  -3.539  0.00044 ***
bs(dis, df = 4, knots = c(4, 7, 11))2 -0.31341    0.01680 -18.658 < 2e-16 ***
bs(dis, df = 4, knots = c(4, 7, 11))3 -0.26618    0.03147  -8.459  3.00e-16 ***
bs(dis, df = 4, knots = c(4, 7, 11))4 -0.39802    0.04647  -8.565 < 2e-16 ***
bs(dis, df = 4, knots = c(4, 7, 11))5 -0.25681    0.09001  -2.853  0.00451 **
bs(dis, df = 4, knots = c(4, 7, 11))6 -0.32926    0.06327  -5.204  2.85e-07 ***
---
Signif. codes:  0 *** 0.001 ** 0.01 * 0.05 . 0.1 1

Residual standard error: 0.06185 on 499 degrees of freedom
Multiple R-squared:  0.7185, Adjusted R-squared:  0.7151
F-statistic: 212.3 on 6 and 499 DF, p-value: < 2.2e-16

> pred = predict(fit, list(dis = grid))
> plot(nox ~ dis, data = Boston, col = "darkgrey")
> lines(grid, pred, col = "red", lwd = 2)
> dev.copy2pdf(file = "MTH522_hw6_p3d.pdf", width = 8, height = 6, out.type = "pdf")
> dev.off()
```





Figure 11: .

After using the `bs()` function to fit a regression spline and predict `nox` using `dis`, the plot shows that all terms in spline fit are significant except `dis` > 10 extreme values.

(e) Now fit a regression spline for a range of degrees of freedom, and plot the resulting fits and report the resulting RSS. Describe the results obtained.

```
> rss = rep(NA, 16)
> for (i in 3:16) {
+   fit = lm(nox ~ bs(dis, df = i), data = Boston)
+   rss[i] = sum(fit$residuals^2) }
> rss[-c(1, 2)]
[1] 1.934107 1.922775 1.840173 1.833966 1.829884 1.816995 1.825653 1.792535 1.796992 1.788999
[11] 1.782350 1.781838 1.782798 1.783546

> which.min(rss)
[1] 14
> rss[which.min(rss)]
[1] 1.781838

> plot(3:16, rss[-c(1, 2)], xlab = "Degrees of freedom", ylab = "RSS", type = "l")
> dev.copy2pdf(file = "MTH522_hw6_p3e.pdf", width = 8, height = 6, out.type = "pdf")
> dev.off()
```




Figure 12: .

From plot, RSS decreases until 14 and then slightly increases after that.

(f) Perform cross-validation or another approach in order to select the best degrees of freedom for a regression spline on this data. Describe your results.

```

> cv = rep(NA, 16)
> for (i in 3:16) {
+   fit = glm(nox ~ bs(dis, df = i), data = Boston)
+   cv[i] = cv.glm(Boston, fit, K = 10)$delta[2]
+ }
> cv[-c(1,2)]
[1] 0.003853321 0.003883472 0.003730361 0.003704349 0.003687914 0.003695486 0.003750861
[8] 0.003680307 0.003681417 0.003699190 0.003745600 0.003771265 0.003719676 0.003761896

```

There were 50 warnings (use warnings() to see them)

```
> warnings()
```

Warning messages:

```

1: In bs(dis, degree = 3L, knots = numeric(0), Boundary.knots = c(1.137, ... :
some 'x' values beyond boundary knots may cause ill-conditioned bases
2: In bs(dis, degree = 3L, knots = numeric(0), Boundary.knots = c(1.137, ... :
some 'x' values beyond boundary knots may cause ill-conditioned bases
3: In bs(dis, degree = 3L, knots = numeric(0), Boundary.knots = c(1.1296, ... :
some 'x' values beyond boundary knots may cause ill-conditioned bases
4: In bs(dis, degree = 3L, knots = numeric(0), Boundary.knots = c(1.1296, ... :
some 'x' values beyond boundary knots may cause ill-conditioned bases
5: In bs(dis, degree = 3L, knots = structure(3.2157, .Names = "50%"), ... :
some 'x' values beyond boundary knots may cause ill-conditioned bases
6: In bs(dis, degree = 3L, knots = structure(3.2157, .Names = "50%"), ... :
some 'x' values beyond boundary knots may cause ill-conditioned bases
7: In bs(dis, degree = 3L, knots = structure(3.1423, .Names = "50%"), ... :
some 'x' values beyond boundary knots may cause ill-conditioned bases
8: In bs(dis, degree = 3L, knots = structure(3.1423, .Names = "50%"), ... :
some 'x' values beyond boundary knots may cause ill-conditioned bases
9: In bs(dis, degree = 3L, knots = structure(c(2.35093333333333, ... :
some 'x' values beyond boundary knots may cause ill-conditioned bases
10: In bs(dis, degree = 3L, knots = structure(c(2.35093333333333, ... :
some 'x' values beyond boundary knots may cause ill-conditioned bases
11: In bs(dis, degree = 3L, knots = structure(c(2.4212, 4.36263333333333 ... :
some 'x' values beyond boundary knots may cause ill-conditioned bases
12: In bs(dis, degree = 3L, knots = structure(c(2.4212, 4.36263333333333 ... :
some 'x' values beyond boundary knots may cause ill-conditioned bases
13: In bs(dis, degree = 3L, knots = structure(c(2.1103, 3.2628, ... :
some 'x' values beyond boundary knots may cause ill-conditioned bases
14: In bs(dis, degree = 3L, knots = structure(c(2.1103, 3.2628, ... :
some 'x' values beyond boundary knots may cause ill-conditioned bases
15: In bs(dis, degree = 3L, knots = structure(c(2.08755, 3.2157, ... :
some 'x' values beyond boundary knots may cause ill-conditioned bases
16: In bs(dis, degree = 3L, knots = structure(c(2.08755, 3.2157, ... :
some 'x' values beyond boundary knots may cause ill-conditioned bases
17: In bs(dis, degree = 3L, knots = structure(c(1.9512, 2.59774, ... :
some 'x' values beyond boundary knots may cause ill-conditioned bases
18: In bs(dis, degree = 3L, knots = structure(c(1.9512, 2.59774, ... :
some 'x' values beyond boundary knots may cause ill-conditioned bases
19: In bs(dis, degree = 3L, knots = structure(c(1.94264, 2.7147, ... :
some 'x' values beyond boundary knots may cause ill-conditioned bases
20: In bs(dis, degree = 3L, knots = structure(c(1.94264, 2.7147, ... :
some 'x' values beyond boundary knots may cause ill-conditioned bases
21: In bs(dis, degree = 3L, knots = structure(c(1.85586666666667, ... :
some 'x' values beyond boundary knots may cause ill-conditioned bases
22: In bs(dis, degree = 3L, knots = structure(c(1.85586666666667, ... :
some 'x' values beyond boundary knots may cause ill-conditioned bases
23: In bs(dis, degree = 3L, knots = structure(c(1.86156666666667, ... :

```

```
> which.min(cv)
[1] 10
> cv[which.min(cv)]
[1] 0.003680307

> plot(3:16, cv[-c(1, 2)], xlab = "Degrees of freedom", ylab = "CV test error", type = "l")
> dev.copy2pdf(file = "MTH522_hw6_p3f.pdf", width = 8, height = 6, out.type = "pdf")
> dev.off()
```




Figure 13: .

CV test error gets minimum at 10 the best degrees of freedom.

Problem 4 (Chapter 7 Exercises 10):

This question relates to the College data set.

(a) Split the data into a training set and a test set. Using out-of-state tuition as the response and the other variables as the predictors, perform forward stepwise selection on the training set in order to identify a satisfactory model that uses just a subset of the predictors.

```
> set.seed(1)
> library(ISLR)
> library(leaps)
> attach(College)
> summary(College)
```

Private	Apps	Accept	Enroll	Top10perc	Top25perc
No :212	Min. : 81	Min. : 72	Min. : 35	Min. : 1.00	Min. : 9.0
Yes:565	1st Qu.: 776	1st Qu.: 604	1st Qu.: 242	1st Qu.:15.00	1st Qu.: 41.0
Median : 1558	Median : 1110	Median : 434	Median :23.00	Median : 54.0	
Mean : 3002	Mean : 2019	Mean : 780	Mean :27.56	Mean : 55.8	
3rd Qu.: 3624	3rd Qu.: 2424	3rd Qu.: 902	3rd Qu.:35.00	3rd Qu.: 69.0	
Max. :48094	Max. :26330	Max. :6392	Max. :96.00	Max. :100.0	
F.Undergrad	P.Undergrad	Outstate	Room.Board	Books	
Min. : 139	Min. : 1.0	Min. : 2340	Min. :1780	Min. : 96.0	
1st Qu.: 992	1st Qu.: 95.0	1st Qu.: 7320	1st Qu.:3597	1st Qu.: 470.0	
Median : 1707	Median : 353.0	Median : 9990	Median :4200	Median : 500.0	
Mean : 3700	Mean : 855.3	Mean :10441	Mean :4358	Mean : 549.4	
3rd Qu.: 4005	3rd Qu.: 967.0	3rd Qu.:12925	3rd Qu.:5050	3rd Qu.: 600.0	
Max. :31643	Max. :21836.0	Max. :21700	Max. :8124	Max. :2340.0	
Personal	PhD	Terminal	S.F.Ratio	perc.alumni	
Min. : 250	Min. : 8.00	Min. : 24.0	Min. : 2.50	Min. : 0.00	
1st Qu.: 850	1st Qu.: 62.00	1st Qu.: 71.0	1st Qu.:11.50	1st Qu.:13.00	
Median :1200	Median : 75.00	Median : 82.0	Median :13.60	Median :21.00	
Mean :1341	Mean : 72.66	Mean : 79.7	Mean :14.09	Mean :22.74	
3rd Qu.:1700	3rd Qu.: 85.00	3rd Qu.: 92.0	3rd Qu.:16.50	3rd Qu.:31.00	
Max. :6800	Max. :103.00	Max. :100.0	Max. :39.80	Max. :64.00	
Expend	Grad.Rate				
Min. : 3186	Min. : 10.00				
1st Qu.: 6751	1st Qu.: 53.00				
Median : 8377	Median : 65.00				
Mean : 9660	Mean : 65.46				
3rd Qu.:10830	3rd Qu.: 78.00				
Max. :56233	Max. :118.00				

```

> train = sample(length(Outstate), length(Outstate)/2)
> test = -train
> College.train = College[train, ]
> College.test = College[test, ]
> fit = regsubsets(Outstate ~ ., data = College.train, nvmax = 17, method = "forward")
> summary = summary(reg.fit)

> par(mfrow = c(1, 3))
> plot(summary$cp, xlab = "Number of Variables", ylab = "Cp", type = "l")
> min.cp = min(summary$cp)
> std.cp = sd(summary$cp)
> abline(h = min.cp + 0.2 * std.cp, col = "red", lty = 2)
> abline(h = min.cp - 0.2 * std.cp, col = "red", lty = 2)

> plot(reg.summary$bic, xlab = "Number of Variables", ylab = "BIC", type = "l")
> min.bic = min(summary$bic)
> std.bic = sd(summary$bic)
> abline(h = min.bic + 0.2 * std.bic, col = "red", lty = 2)
> abline(h = min.bic - 0.2 * std.bic, col = "red", lty = 2)

> plot(summary$adjr2, xlab = "Number of Variables", ylab = "Adjusted R2",
+       type = "l", ylim = c(0.4, 0.84))
> max.adj2 = max(summary$adjr2)
> std.adj2 = sd(summary$adjr2)
> abline(h = max.adj2 + 0.2 * std.adj2, col = "red", lty = 2)
> abline(h = max.adj2 - 0.2 * std.adj2, col = "red", lty = 2)
> dev.copy2pdf(file = "MTH522_hw6_p4a.pdf", width = 8, height = 6, out.type = "pdf")
> dev.off()

```




Figure 14: .

We are trying to find the minimum size for the subset for which the scores are within 0.2 standard deviations of optimum. From the plot, Cp, BIC and adjr2 show that size should be bigger than 5, so we may see that number of variables 6 is the minimum size. Let's see the best 6 variables;

```

> fit = regsubsets(Outstate ~ ., data = College, method = "forward")
> coefi = coef(fit, id = 6)
> names(coefi)
[1] "(Intercept)" "PrivateYes" "Room.Board" "PhD" "perc.alumni" "Expend"
[7] "Grad.Rate"

```

(b) Fit a GAM on the training data, using out-of-state tuition as the response and the features selected in the previous step as the predictors. Plot the results, and explain your findings.

```

> fit = gam(Outstate ~ Private + s(Room.Board, df = 2) + s(PhD, df = 2) +
+ s(perc.alumni, df = 2) + s(Expend, df = 5) + s(Grad.Rate, df = 2), data = College.train)
> par(mfrow = c(2, 3))
> plot(fit, se = T, col = "blue")
> dev.copy2pdf(file = "MTH522_hw6_p4b.pdf", width = 8, height = 6, out.type = "pdf")
> dev.off()

```

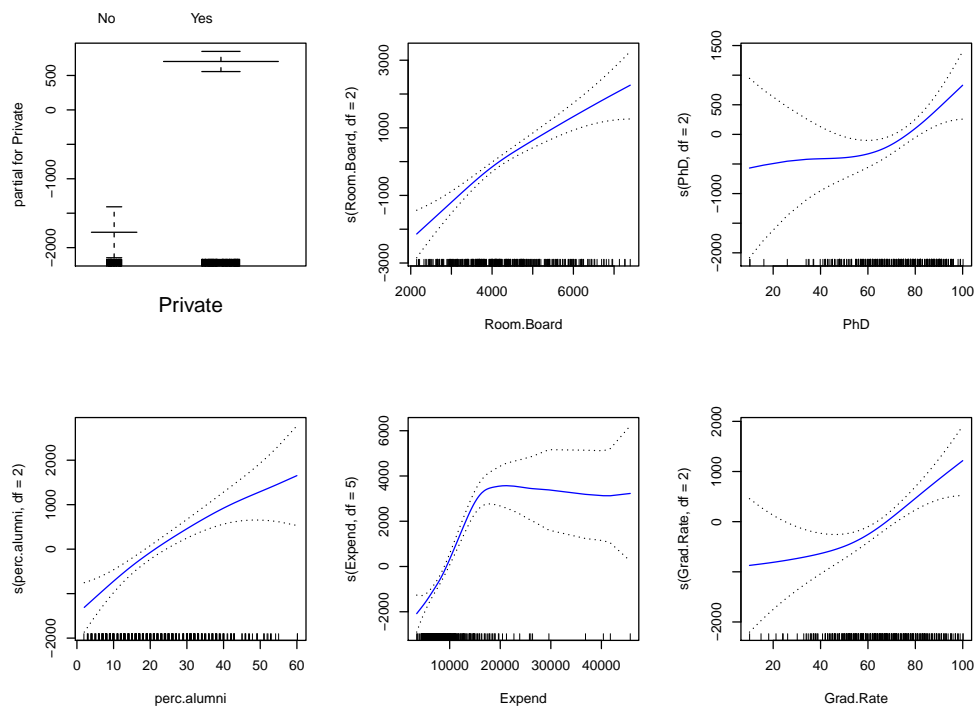


Figure 15: .

(c) Evaluate the model obtained on the test set, and explain the results obtained.

```
> pred = predict(fit, College.test)
> err = mean((College.test$Outstate - pred)^2)
> err
[1] 3745460

> tss = mean((College.test$Outstate - mean(College.test$Outstate))^2)
> rss = 1 - err/tss
> rss
[1] 0.7696916
```

After evaluation the model obtained on the test set, we obtain R^2 of 0.7696916 using GAM with 6 predictors.

(d) For which variables, if any, is there evidence of a non-linear relationship with the response?

```
> summary(fit)

Call: gam(formula = Outstate ~ Private + s(Room.Board, df = 2) + s(PhD,
df = 2) + s(perc.alumni, df = 2) + s(Expend, df = 5) + s(Grad.Rate,
df = 2), data = College.train)
Deviance Residuals:
Min       1Q   Median       3Q      Max
-4977.74 -1184.52   58.33  1220.04  7688.30

(Dispersion Parameter for gaussian family taken to be 3300711)

Null Deviance: 6221998532 on 387 degrees of freedom
Residual Deviance: 1231165118 on 373 degrees of freedom
AIC: 6941.542

Number of Local Scoring Iterations: 2

Anova for Parametric Effects
Df      Sum Sq    Mean Sq F value    Pr(>F)
Private      1 1779433688 1779433688 539.106 < 2.2e-16 ***
s(Room.Board, df = 2) 1 1221825562 1221825562 370.171 < 2.2e-16 ***
s(PhD, df = 2)      1  382472137  382472137 115.876 < 2.2e-16 ***
s(perc.alumni, df = 2) 1  328493313  328493313  99.522 < 2.2e-16 ***
s(Expend, df = 5)    1  416585875  416585875 126.211 < 2.2e-16 ***
s(Grad.Rate, df = 2) 1   55284580   55284580  16.749 5.232e-05 ***
Residuals      373 1231165118   3300711
---
Signif. codes:  0 *** 0.001 ** 0.01 * 0.05 . 0.1 1

Anova for Nonparametric Effects
Npar Df  Npar F      Pr(F)
(Intercept)
Private
s(Room.Board, df = 2)      1  3.5562  0.06010 .
s(PhD, df = 2)            1  4.3421  0.03786 *
s(perc.alumni, df = 2)    1  1.9158  0.16715
s(Expend, df = 5)         4 16.8636 1.016e-12 ***
s(Grad.Rate, df = 2)      1  3.7208  0.05450 .
---
Signif. codes:  0 *** 0.001 ** 0.01 * 0.05 . 0.1 1
```

Anova for Nonparametric shows that there is a strong evidence of non-linear relationship between "Outstate" and "Expend". Using p value of 0.05, there is a moderately strong non-linear relationship between "Outstate" and "PhD" or "Grad.Rate".